

Contents

| | |
|---|----------|
| 1 Chapter 06 — Shortest Paths | 1 |
| 1.1 6.1 Problem definition | 1 |
| 1.2 6.2 Relaxation and the upper-bound invariant | 1 |
| 1.2.1 Definition (Relaxation) | 1 |
| 1.2.2 Invariant 6.1 (Upper-bound invariant) | 2 |
| 1.2.3 Lemma 6.2 (Witness paths via predecessors) | 2 |
| 1.3 6.3 Dijkstra (nonnegative weights) | 2 |
| 1.3.1 Algorithm | 2 |
| 1.3.2 Theorem 6.3 (Finalization lemma) | 2 |
| 1.3.3 Corollary 6.4 | 2 |
| 1.3.4 Complexity | 2 |
| 1.4 6.4 Bellman–Ford (negative edges allowed) | 2 |
| 1.4.1 Algorithm | 2 |
| 1.4.2 Lemma 6.5 (k-edge optimality) | 3 |
| 1.4.3 Theorem 6.6 (Correctness without negative cycles) | 3 |
| 1.4.4 Theorem 6.7 (Negative-cycle detection) | 3 |
| 1.5 Appendix — Trace events as proof witnesses | 3 |

% Chapter 06 — Shortest Paths (Dijkstra & Bellman–Ford) % CS161 Reader (Plotkin F10) %

1 Chapter 06 — Shortest Paths

Shortest-path algorithms are relaxation machines: they maintain **upper bounds** and improve them until no improvement is possible. Dijkstra is fast but needs **nonnegative weights**; Bellman–Ford is slower but handles negative edges and detects negative cycles.

Primary patterns: relaxation, invariant.

Executable witnesses: `cs161lab.algorithms.sssp.dijkstra`, `cs161lab.algorithms.sssp.bellman_ford`.

1.1 6.1 Problem definition

Let $G = (V, E)$ be a directed graph with weights $w : E \rightarrow \mathbb{R}$ and a source s . Define

$$\delta(s, v) = \min_{p: s \leadsto v} w(p),$$

where $w(p)$ sums edge weights along p ; if no path exists, $\delta(s, v) = +\infty$.

If there is a reachable negative cycle, shortest paths are ill-defined for vertices reachable from it (distance $-\infty$).

We maintain estimates $d[v]$ and predecessors $\pi[v]$.

1.2 6.2 Relaxation and the upper-bound invariant

1.2.1 Definition (Relaxation)

Relaxing edge (u, v) sets

$$d[v] \leftarrow \min\{d[v], d[u] + w(u, v)\}.$$

1.2.2 Invariant 6.1 (Upper-bound invariant)

At all times,

$$d[v] \geq \delta(s, v) \quad \forall v.$$

Proof. True at initialization. If $d[u] \geq \delta(s, u)$ then $d[u] + w(u, v) \geq \delta(s, u) + w(u, v) \geq \delta(s, v)$, so taking a min preserves the inequality. \square

1.2.3 Lemma 6.2 (Witness paths via predecessors)

Whenever $d[v]$ is finite, predecessors reconstruct an $s \rightarrow v$ path of weight $d[v]$.

Proof. Set $\pi[v] = u$ exactly when relaxation improves $d[v]$ using u . Following predecessors yields a path. \square

1.3 6.3 Dijkstra (nonnegative weights)

Assume $w(e) \geq 0$ for all edges.

1.3.1 Algorithm

Initialize $d[s] = 0$, others $+\infty$. Maintain finalized set $S = \emptyset$. Repeatedly extract $u \notin S$ with minimum $d[u]$, add to S , relax outgoing edges.

Trace events: `finalize`, `relax`.

1.3.2 Theorem 6.3 (Finalization lemma)

When Dijkstra finalizes a vertex u , then $d[u] = \delta(s, u)$.

Proof. Let S be finalized vertices before choosing u . Consider a shortest path from s to u and let (x, y) be the first edge crossing from S to $V \setminus S$. When x was finalized, (x, y) was relaxed, so

$$d[y] \leq d[x] + w(x, y) = \delta(s, x) + w(x, y) = \delta(s, y).$$

With Invariant 6.1, $d[y] = \delta(s, y)$. Since remaining edges on the path are nonnegative, $\delta(s, u) \geq \delta(s, y) = d[y]$. As u minimizes d outside S , $d[u] \leq d[y] \leq \delta(s, u)$. With Invariant 6.1, $d[u] \geq \delta(s, u)$, hence equality. \square

1.3.3 Corollary 6.4

At termination, Dijkstra outputs correct distances for all vertices reachable from s .

1.3.4 Complexity

Using a binary heap: $O((n + m) \log n)$.

1.4 6.4 Bellman–Ford (negative edges allowed)

1.4.1 Algorithm

Initialize $d[s] = 0$, others $+\infty$. Repeat $n - 1$ passes: relax every edge. Do one extra pass; if any edge relaxes, report a reachable negative cycle.

Trace events: `iteration`, `relax`, `neg_cycle`.

1.4.2 Lemma 6.5 (k-edge optimality)

After k full passes, $d[v]$ is at most the shortest-path weight among all $s \rightarrow v$ paths using at most k edges.

Proof (induction). Each pass allows paths that extend a k -edge best path by one edge. \square

1.4.3 Theorem 6.6 (Correctness without negative cycles)

If there is no reachable negative cycle, then after $n - 1$ passes, $d[v] = \delta(s, v)$ for all v .

Proof. A shortest path can be chosen simple, hence has at most $n - 1$ edges. Apply Lemma 6.5 with $k = n - 1$ and combine with Invariant 6.1. \square

1.4.4 Theorem 6.7 (Negative-cycle detection)

If an edge can still be relaxed after $n - 1$ passes, then a reachable negative cycle exists.

Proof. A strict improvement implies an improving walk of length at least n edges, hence repeats a vertex; the repeated segment is a negative cycle. \square

1.5 Appendix — Trace events as proof witnesses

- Dijkstra `finalize` events correspond to Theorem 6.3 checkpoints.
- Bellman–Ford `iteration(k)` tracks the Lemma 6.5 induction parameter.
- `neg_cycle` is a concrete witness of a violated “no negative cycle” assumption.